

Class 6: Interpreting Logistic Regression

```
### Load packages:
#library(knitr)
library(tidyverse)
#library(data.table)
#library(janitor)
#library(haven)
#library(readxl)
#library(psych)
#library(statar)
library(mktg482)
library(sjPlot)
```

We often want to interpret variables in a model. Let's try this on the Tuango RFM dataset.

```
load("../Data/Tuango_rfm.Rdata")
```

First, we estimate our logistic regression model:

```
lr <- glm(buyer ~ recency + frequency + monetary + platform + category,
          family = binomial, data = tuango)
```

As we discussed, logistic regression coefficient estimates are hard to interpret:

```
summary(lr)
```

Call:

```
glm(formula = buyer ~ recency + frequency + monetary + platform +
     category, family = binomial, data = tuango)
```

Deviance Residuals:

	Min	1Q	Median	3Q	Max
	-1.8188	-0.2831	-0.2518	-0.1785	4.1222

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-3.4904692	0.1402524	-24.887	< 2e-16	***
recency	-0.0156296	0.0019678	-7.943	1.98e-15	***
frequency	0.1186148	0.0201425	5.889	3.89e-09	***
monetary	0.0012814	0.0002386	5.370	7.88e-08	***
platformBrowser	-0.3151997	0.2390529	-1.319	0.187	
category	0.0153914	0.0121877	1.263	0.207	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


(Dispersion parameter for binomial family taken to be 1)

Null deviance:	3852.0	on 13938	degrees of freedom
Residual deviance:	3651.4	on 13933	degrees of freedom
AIC:	3663.4		

Number of Fisher Scoring iterations: 8

The `sjPlot` package helps assess the effects of predictor variables, in particular via the `plot_model` command. We can use the package to display the coefficients together with confidence intervals. Make sure you use the options `show.values = TRUE`, `transform = NULL`.

```
plot_model(lr, show.values = TRUE, transform = NULL)
```



Nonetheless, it is still difficult to interpret what these estimates mean. Instead, we use two approaches to more easily interpret logistic regression coefficient estimates: “Variable Importance” and “Marginal Effects Plots”

Variable Importance

Often we want to know how important a given variable is relative to other variables. One common way to do this is to normalize the coefficients by the standard deviation of the variable. A problem with this idea is that coefficient estimates for continuous and binary variables not comparable. Andrew Gelman proposed a solution in, “Scaling regression inputs by dividing by two standard deviations,” *Statistics in Medicine* (2008), Vol. 27, pp. 2965-2873.

The idea is to measure the effect of continuous variables by considering the effect of a 2 standard deviation change in the variable, and leaving the effect of binary variables unchanged. Since the SD of an evenly-split binary variable (i.e., one with an equal number of 0s and 1s) is 0.5, a 0->1 “one unit change” in a binary variable is equivalent to a 2 SD change. Consequently, the coefficient represents the effect a 2*SD change for such variables (in addition to the convention interpretation as the effect of a one unit change), and now we can compare across the two kinds of variables.

This idea is implemented in R functions I have created for you: `varimp.logistic()` and its companion plotting function `plotimp.logistic()` which can be piped. To use these functions, you need to use load the `mktg482` library as we did using the command `library(mktg482)`; if that command fails for you, you need to install this package by typing:

```
devtools::install_github("fzettelmeyer/mktg482", upgrade = "never", force = TRUE)
```

into the console (you only need to do this once).

The syntax for using the `varimp.logistic()` and `plotimp.logistic()` functions is:

```
varimp.logistic(lrmodel)
```

```
plotimp.logistic(lrmodel)
```


Let's try these functions on the Tuango RFM dataset. The first gives a table that lists each variable, whether it is a factor or not, its variable importance before taking the absolute value (so as to preserve information about the sign of the effect), a confidence interval for the variable importance, and the p-value from the logistic regression. The second is a figure which shows the absolute value of variable importance, where information on the sign of the effect is provided via the color of the x-axis label (black=positive; red=negative).

```
varimp.logistic(lr)
```

```
# A tibble: 5 x 6
```

	variable	factor	var_imp	var_imp_lower	var_imp_upper	p_value
	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	recency	No	-2.41	-3.01	-1.82	0
2	frequency	No	0.504	0.337	0.672	0
3	platformBrowser	Yes	-0.315	-0.784	0.153	0.187
4	monetary	No	0.298	0.189	0.406	0
5	category	No	0.124	-0.0685	0.317	0.207

```
varimp.logistic(lr) %>% plotimp.logistic()
```



```
# A tibble: 5 x 6
```


	variable	factor	var_imp	var_imp_lower	var_imp_upper	p_value
	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	recency	No	-2.41	-3.01	-1.82	0
2	frequency	No	0.504	0.337	0.672	0
3	platformBrowser	Yes	-0.315	-0.784	0.153	0.187
4	monetary	No	0.298	0.189	0.406	0
5	category	No	0.124	-0.0685	0.317	0.207

Marginal Effects

Another way to interpret the coefficient of a variable is to analyze the “marginal effect” of the variable. This corresponds to asking “Holding all other variables at their current values, what is the effect of changing Variable X on the model prediction?” We can get this by using the `plot_model` command with the `type = "eff"` option. This shows the effect of each variable in turn:

```
plot_model(lr, type= "eff")
```

\$recency




\$frequency



\$monetary



\$platform



\$category



If you want to look at one variable only, use the 'terms = "FEATURE"' option where FEATURE is the variable for which you want to see the marginal effect.

```
plot_model(lr, type= "eff", terms = "recency")
```

