



CREATIVE GAMING

Assignment for

Creative Gaming: Uplift Modeling

You decide to try uplift modeling to pick consumers for Zalon. You plan to use the experimental control group ("**Group 1**"), i.e., the 30,000 customers who did not receive an ad), as well as "**Group 2**", i.e., the randomly chosen experimental group of 30,000 customers who did receive the ad.

Part 1: Uplift Modeling Using Machine Learning (Random Forests)

1. Prepare your data:

Hint: Please visualize what you are doing by looking at the new data frames you create in each step.

- a) Create "Group 2" by sampling from `cg_ad_treatment`. You can use this syntax:

```
cg_ad_random <- cg_ad_treatment[sample_random_30000,]
```

- b) Create a stacked dataset for the uplift analysis by combining `cg_organic_control` (Group 1) and `cg_ad_random` (Group 2). You can use this syntax:

```
expdata_stacked <-  
  rbind(cg_organic_control %>% mutate(ad = 0),  
        cg_ad_random %>% mutate(ad = 1))
```

- c) Split the stacked dataset into a training and test dataset. You can use this syntax:

```
set.seed(1234)  
split.index <- stratified(expdata_stacked, c("ad",  
                                             "converted"), 0.7, bothSets=TRUE)  
expdata_stacked.train <- split.index[[1]]  
expdata_stacked.test <- split.index[[2]]
```

Professor Florian Zettelmeyer prepared this case to provide material for class discussion rather than to illustrate either effective or ineffective handling of a business situation. Events and people in this case are fictionalized.

2. Train an uplift model using random forests. Add the predicted scores for the treatment and control models to `expdata_stacked.test` and calculate the uplift score.

Hint: Please see the handout “Random Forests in R.pdf” for how to run a random forest.

3. Calculate the Uplift (%) and Incremental Uplift (%) for the uplift model (use 20 instead of the standard 10 groups) and plot performance metrics. Interpret the plots.
4. Using the `incremental_Y1` column from the performance metric table created by `QiniTable()`, calculate the incremental profit you expect to make if you targeted the best 30,000 consumers of 120,000 using the uplift model.

Hint: For every n-tile, the `incremental_Y1` tells you how many incremental purchases were made when consumers up to that n-tile were targeted. To extrapolate correctly to picking the best 30,000 from 120,000, notice that there are a total of 9,000 consumers who got the ad in the test sample `expdata_stacked.test`.

5. Calculate the Uplift (%) and Incremental Uplift (%) you would get if you used a propensity model (use 20 instead of the standard 10 groups). Compare the Uplift (%) performance metric between the uplift and propensity models. Interpret the difference.

Hint: To compare the performance of the uplift and propensity models, use the functions `QiniCurve()` and `QiniBarPlot()`

6. Using the `incremental_Y1` column from the performance metric table created by `QiniTable()` for the propensity model, calculate the incremental profit you expect to make if you targeted the best 30,000 consumers of 120,000 using the propensity model. How much more money do you expect to make from using an uplift instead of a propensity model?

Part 2: Targeting the optimal percent of customers

So far we have always targeted a 25% of model-selected customers (by picking the best 30,000 out of 120,000 customers). We now want to evaluate whether we should target fewer or more than 25% of customers.

1. What formula would you use to select which consumers to target using a propensity model where your goal is to maximize profits? What percentage of customers in the ad treatment group of `expdata_stacked.test` would you target using the propensity model?
2. What formula would you use to select which consumers to target using an uplift model where your goal is to maximize incremental profits. What percentage of

customers in the ad treatment group of `expdata_stacked.test` would you target using the uplift model?

3. Rounding the targeting percentage numbers you calculated in 1. and 2. to the nearest 5%, use the `QiniTable()` you calculated for the propensity and uplift models in Part 1 to calculate the incremental profits you would have obtained in the `expdata_stacked.test` dataset if you had targeted the optimal percentage of customers suggested by each model.
4. Give two reasons for why one model beats the other in incremental profits.